

Отчет по домашнему заданию №2 курса Highload Architect

Содержание отчета

Mysql.....	2
DDL Таблицы users:.....	2
До добавления индекса.....	3
Explane запроса:.....	3
После добавления индекса.....	3
Добавление индекса:.....	3
Explane запроса:.....	3
Таблицы и графики.....	3
Выводы.....	4
Вопросы:.....	4
Приложение 1 – лог запросов до добавления индекса.....	4
Приложение 2 – лог запросов после добавления индекса.....	8
Postgres.....	12
DDL Таблицы users:.....	12
До добавления индекса.....	12
После добавления индекса.....	13
Добавление индекса:.....	13
Explane запроса:.....	13
Таблицы и графики.....	13
Выводы.....	14
Вопросы:.....	14
Приложение 1 – лог запросов до добавления индекса.....	14
Приложение 2 – лог запросов после добавления индекса.....	18

Mysql

DDL Таблицы users:

```
SHOW CREATE TABLE users;
```

```
| Table | Create Table
|
+-----+
| users | CREATE TABLE `users` (
  `user_id` char(36) NOT NULL,
  `first_name` varchar(64) NOT NULL,
  `second_name` varchar(64) DEFAULT NULL,
  `sex` enum('male','female') DEFAULT NULL,
  `biography` text,
  `city` varchar(64) DEFAULT NULL,
  `birthdate` date DEFAULT NULL,
  KEY `users_user_id_index` (`user_id`),
  KEY `users_first_name_second_name_index` (`first_name`,`second_name`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci |
+-----+
+-----1 row in set (0.00 sec)
```

```
mysql> SHOW INDEX FROM users;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible	Expression
users	1	users_user_id_index	1	user_id	A	922210		NULL	NULL		BTREE		YES	NULL
users	1	users_first_name_second_name_index	1	first_name	A	139		NULL	NULL		BTREE		YES	NULL
users	1	users_first_name_second_name_index	2	second_name	A	57115		NULL	NULL	YES	BTREE		YES	NULL

3 rows in set (0.01 sec)

До добавления индекса

Explane запроса:

mysql> EXPLAIN select * from users where first_name like 'Иван%' and second_name like 'Бес%';

```
+-----+
| EXPLAIN |
+-----+
| -> Filter: ((users.first_name like 'Иван%') and (users.second_name like 'Бес%')) (cost=107914 rows=12264)
  -> Table scan on users (cost=107914 rows=993595)
    |
+-----+
```

1 row in set (0.00 sec)

После добавления индекса

Добавление индекса:

mysql> create index users_first_name_second_name_index
-> on users (first_name, second_name);
Query OK, 0 rows affected (3.05 sec)
Records: 0 Duplicates: 0 Warnings: 0

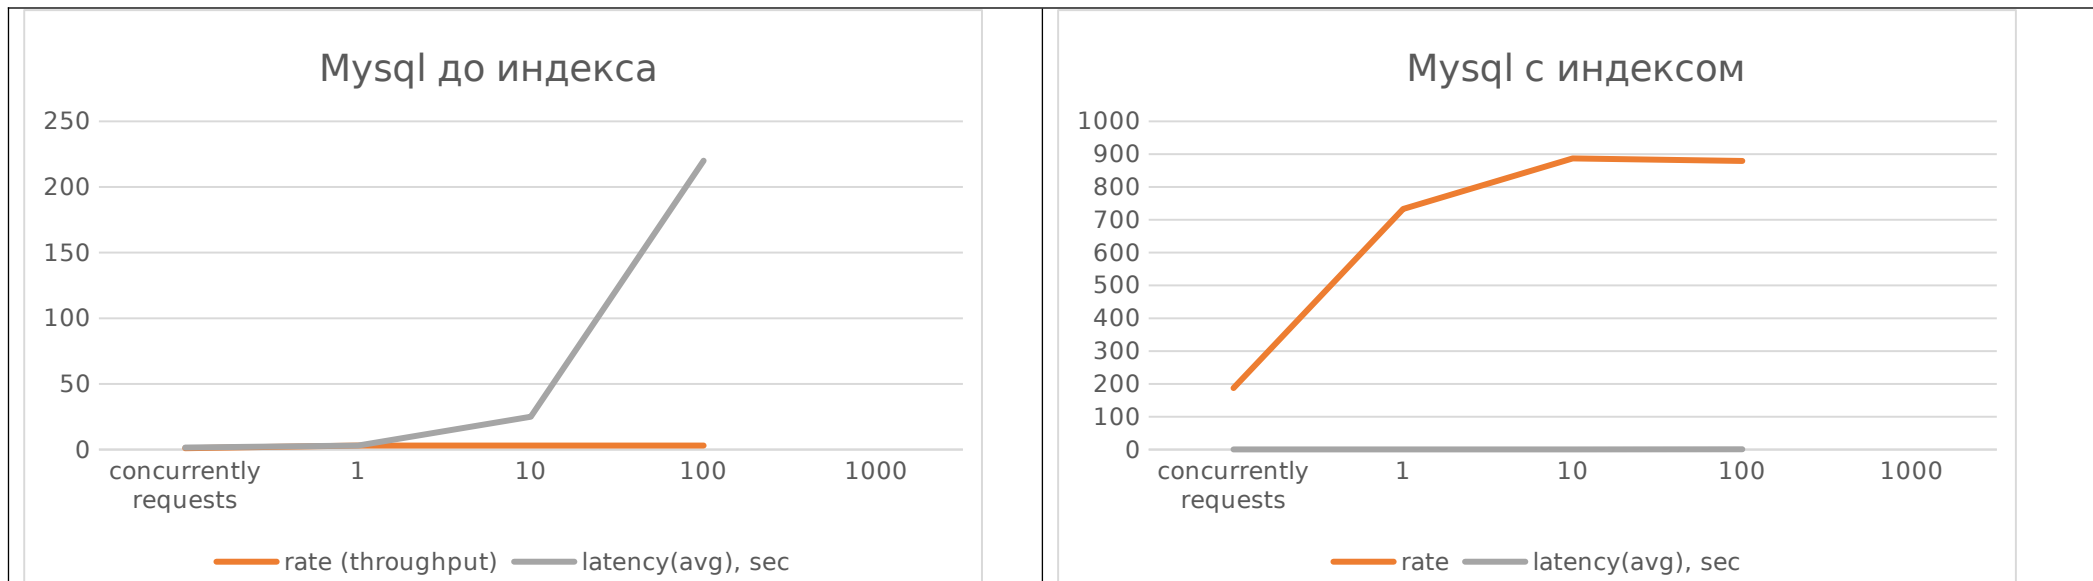
Explane запроса:

```
mysql> EXPLAIN select * from users where first_name like 'Иван%' and second_name like 'Бес%';
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | users | NULL | range | users_first_name_second_name_index | users_first_name_second_name_index | 517 | NULL | 37052 | 11.11 | Using index condition; Using MRR |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

Таблицы и графики

Mysql					
Without index			With index		
concurrently requests	rate (throughput)	latency(avg), sec	concurrently requests	rate	latency(avg), sec

1	1	1,5	1	187	0,005
10	3	3	10	733	0,012
100	3	25	100	887	0,08
1000	3	220	1000	879	0,76



Выводы

1. Поведение после добавления индекса ожидаемо, explain показывает использование индекса
2. Без использования индекса throughput составляет всего 3 запроса в секунду. После добавления индекса сервер без деградации latency держит примерно до 800 запросов в секунду. После наблюдается увеличение времени отклика.

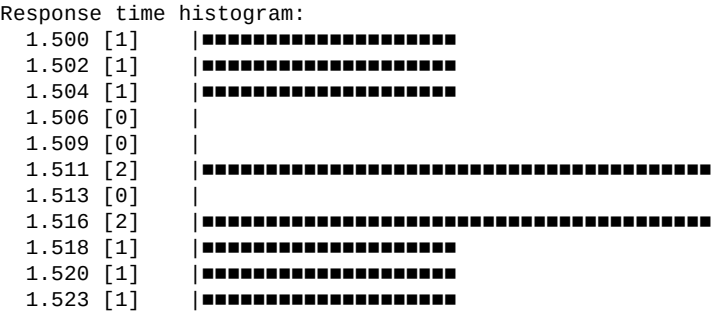
Вопросы:

1. Как можно увеличить throughput сервера. Вероятно, какими-то настройками, но какими?

Приложение 1 – лог запросов до добавления индекса

```
hey -n 10 -c 1 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'
```

Summary:
Total: 15.1099 secs
Slowest: 1.5226 secs
Fastest: 1.4996 secs
Average: 1.5110 secs
Requests/sec: 0.6618



Latency distribution:
10% in 1.4997 secs
25% in 1.5100 secs
50% in 1.5138 secs
75% in 1.5186 secs
90% in 1.5226 secs
0% in 0.0000 secs
0% in 0.0000 secs

Details (average, fastest, slowest):
DNS+ dialup: 0.0000 secs, 1.4996 secs, 1.5226 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0001 secs
req write: 0.0000 secs, 0.0000 secs, 0.0001 secs
resp wait: 1.5108 secs, 1.4995 secs, 1.5225 secs
resp read: 0.0001 secs, 0.0001 secs, 0.0001 secs

Status code distribution:
[200] 10 responses

```
hey -n 100 -c 10 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек' >>mysql.log
```

Summary:
Total: 32.6865 secs
Slowest: 6.5508 secs
Fastest: 1.6266 secs
Average: 3.0561 secs
Requests/sec: 3.0594

Response time histogram:

Value	Index	Frequency
1.627	[1]	1
2.119	[34]	34
2.611	[0]	1
3.104	[0]	1
3.596	[48]	48
4.089	[0]	1
4.581	[0]	1
5.074	[12]	12
5.566	[0]	1
6.058	[0]	1
6.551	[5]	5

Latency distribution:

```
10% in 1.6340 secs
25% in 1.6385 secs
50% in 3.2664 secs
75% in 3.2776 secs
90% in 4.9008 secs
95% in 6.5221 secs
99% in 6.5508 secs
```

Details (average, fastest, slowest):

```
DNS+dialup: 0.0001 secs, 1.6266 secs, 6.5508 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0004 secs
req write: 0.0000 secs, 0.0000 secs, 0.0003 secs
resp wait: 3.0560 secs, 1.6265 secs, 6.5501 secs
resp read: 0.0000 secs, 0.0000 secs, 0.0001 secs
```

Status code distribution:

[200] 100 responses

```
hey -n 1000 -c 100 -t 0 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бес' >>mysql.log
```

Summary:

```
Total:      337.3888  secs
Slowest:    297.5786  secs
Fastest:    1.5543   secs
Average:    25.2815   secs
Requests/sec: 2.9639
```

Response time histogram:

```

1.554 [1] | 
31.157 [779] | ████████████████████████████████████████████████████████
60.759 [121] | ████████
90.362 [40] | ███
119.964 [20] | █
149.566 [11] | █
179.169 [11] | █
208.771 [7] | 

```

```
238.374 [1] |
267.976 [6] |
297.579 [3] |
```

Latency distribution:

```
10% in 3.3414 secs
25% in 5.0195 secs
50% in 11.4449 secs
75% in 26.8221 secs
90% in 60.3494 secs
95% in 106.0476 secs
99% in 234.8641 secs
```

Details (average, fastest, slowest):

```
DNS+dialup: 0.0006 secs, 1.5543 secs, 297.5786 secs
DNS-lookup: 0.0001 secs, 0.0000 secs, 0.0088 secs
req write: 0.0001 secs, 0.0000 secs, 0.0068 secs
resp wait: 25.2806 secs, 1.5542 secs, 297.5786 secs
resp read: 0.0000 secs, 0.0000 secs, 0.0004 secs
```

Status code distribution:

[200] 1000 responses

```
hey -n 3000 -c 1000 -t 0 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бес'
```

Summary:

```
Total:      987.2750  secs
Slowest:    982.3998  secs
Fastest:    1.5457   secs
Average:    220.2241  secs
Requests/sec: 3.0387
```

Response time histogram:

[illegible]

Latency distribution:

```
10% in 10.0489 secs
25% in 43.0538 secs
50% in 141.9484 secs
```

```
Details (average, fastest, slowest):
DNS+dialup: 0.0209 secs, 1.5457 secs, 982.3998 secs
DNS-lookup: 0.0066 secs, 0.0000 secs, 0.1053 secs
req write: 0.0068 secs, 0.0000 secs, 0.0609 secs
resp wait: 220.1956 secs, 1.5456 secs, 982.2938 secs
resp read: 0.0000 secs, 0.0000 secs, 0.0003 secs

Status code distribution:
[200] 3000 responses
```

```
hey -n 10 -c 1 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'
```

```
Total:      0.0534  secs
Slowest:    0.0064  secs
Fastest:    0.0049  secs
Average:    0.0053  secs
Requests/sec: 187.4118
```

[illegible]

```
10% in 0.0049 secs
25% in 0.0050 secs
50% in 0.0051 secs
75% in 0.0063 secs
90% in 0.0064 secs
0% in 0.0000 secs
0% in 0.0000 secs
```

```
DNS+dialup: 0.0000 secs, 0.0049 secs, 0.0064 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0002 secs
req write: 0.0000 secs, 0.0000 secs, 0.0000 secs
```



```
hey -n 1000 -c 100 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бес'
```

Summary:

```
Total:      1.1261  secs
Slowest:    0.9258  secs
Fastest:    0.0050  secs
Average:    0.0839  secs
Requests/sec: 887.9856
```

Response time histogram:

[illegible]

Latency distribution:

```
10% in 0.0087 secs
25% in 0.0139 secs
50% in 0.0323 secs
75% in 0.0866 secs
90% in 0.2173 secs
95% in 0.3598 secs
99% in 0.7354 secs
```

Details (average, fastest, slowest):

```
DNS+ dialup: 0.0002 secs, 0.0050 secs, 0.9258 secs
DNS- lookup: 0.0000 secs, 0.0000 secs, 0.0019 secs
req write: 0.0001 secs, 0.0000 secs, 0.0027 secs
resp wait: 0.0836 secs, 0.0049 secs, 0.9257 secs
resp read: 0.0000 secs, 0.0000 secs, 0.0003 secs
```

Status code distribution:

[200] 1000 responses

```
hey -n 3000 -c 1000 -t 0 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бес'
```

Summary:

```
Total:      3.4099  secs
Slowest:    3.3577  secs
Fastest:    0.0049  secs
Average:    0.7616  secs
Requests/sec: 879.7935
```

Response time histogram:

[illegible]

0.675	[558]	■■■■■■■■■■■■■■■■■■■■■■■
1.011	[393]	■■■■■■■■■■■■■■■■■■■
1.346	[267]	■■■■■■■■■
1.681	[203]	■■■■■■■
2.017	[137]	■■■■■
2.352	[93]	■■■
2.687	[73]	■■
3.022	[58]	■■
3.358	[31]	■

Latency distribution:

```
10% in 0.0422 secs
25% in 0.1712 secs
50% in 0.5103 secs
75% in 1.1308 secs
90% in 1.9137 secs
95% in 2.4300 secs
99% in 3.0357 secs
```

Details (average, fastest, slowest):

```
DNS+dialup: 0.0093 secs, 0.0049 secs, 3.3577 secs
DNS-lookup: 0.0062 secs, 0.0000 secs, 0.0862 secs
req write: 0.0014 secs, 0.0000 secs, 0.0529 secs
resp wait: 0.7445 secs, 0.0049 secs, 3.3091 secs
resp read: 0.0003 secs, 0.0000 secs, 0.0522 secs
```

Status code distribution:

[200] 3000 responses

Postgres

DDL Таблицы users:

```
\d users
```

Table "public.users"

Column	Type	Collation	Nullable	Default
user_id	character(36)		not null	
first_name	character varying(64)		not null	
second_name	character varying(64)			
sex	sex_status			
biography	text			
city	character varying(64)			
birthdate	timestamp without time zone			

Indexes:

"users_pkey" PRIMARY KEY, btree (user_id)

"users_first_name_second_name_index" btree (first_name, second_name)

Referenced by:

TABLE "token" CONSTRAINT "tokens_user_id_fkey" FOREIGN KEY (user_id) REFERENCES users(user_id) ON DELETE CASCADE

TABLE "user_credentials" CONSTRAINT "user_credentials_user_id_fkey" FOREIGN KEY (user_id) REFERENCES users(user_id) ON DELETE CASCADE

До добавления индекса

```
snet=# EXPLAIN select * from users where first_name like 'Иван%' and second_name like 'Бес%';  
          QUERY PLAN
```

```
-----  
Gather  (cost=1000.00..22590.04 rows=133 width=128)
```

```
Workers Planned: 2
```

```
->  Parallel Seq Scan on users  (cost=0.00..21576.74 rows=55 width=128)
```

Filter: (((first_name)::text ~~ 'Иван% '::text) AND ((second_name)::text ~~ 'Бес% '::text))
(4 rows)

После добавления индекса

Добавление индекса:

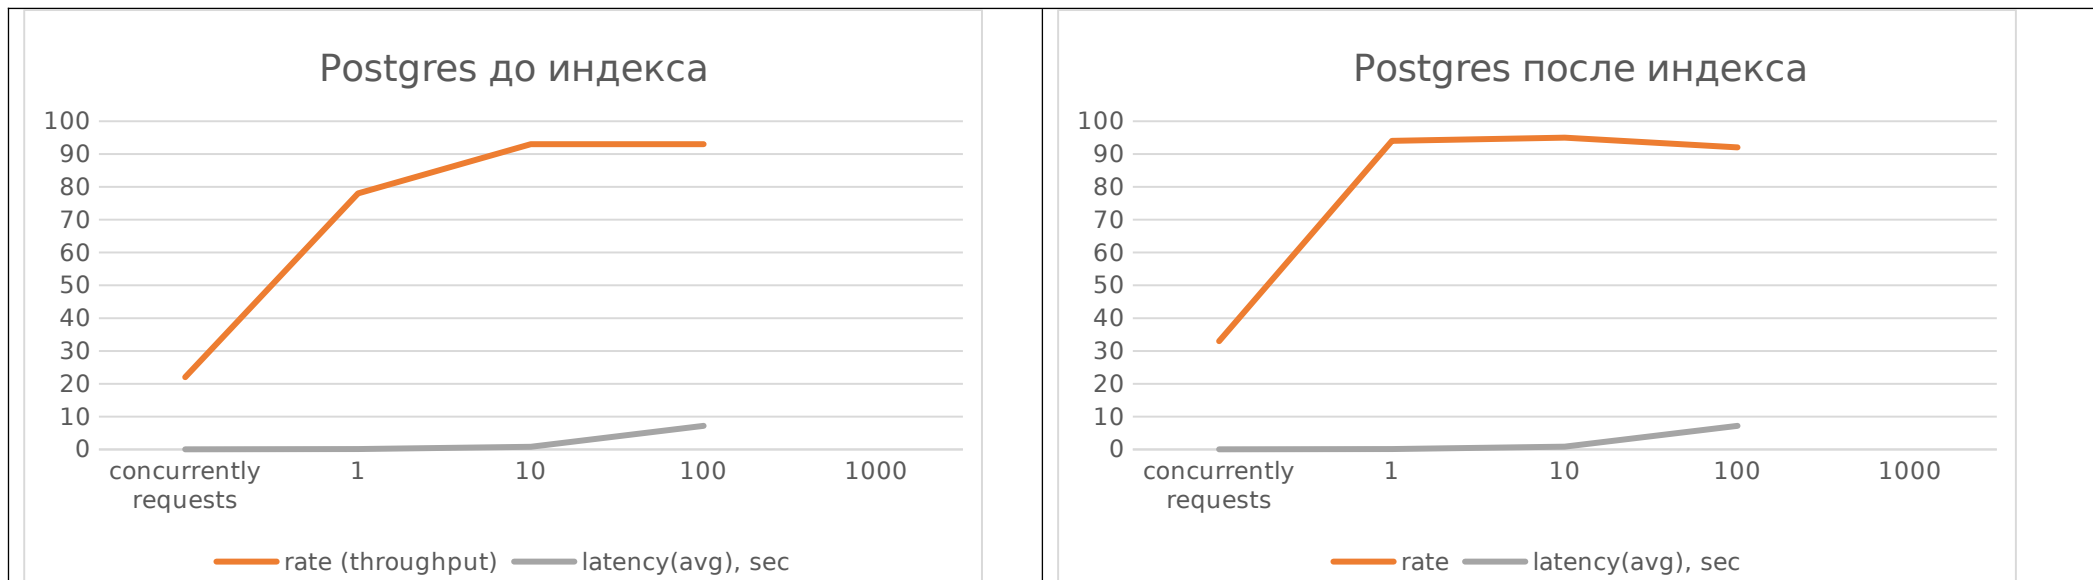
```
snat=# create index users_first_name_second_name_index
      on public.users (first_name, second_name);
CREATE INDEX
```

Explane запроса:

```
EXPLAIN analyze select first_name from users where first_name like 'Иван%' and second_name like 'Бес%';
              QUERY PLAN
-----
Gather  (cost=1000.00..22578.10 rows=1 width=13) (actual time=24.320..29.112 rows=0 loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on users  (cost=0.00..21578.00 rows=1 width=13) (actual time=22.741..22.743 rows=0 loops=3)
        Filter: (((first_name)::text ~~ 'Иван% '::text) AND ((second_name)::text ~~ 'Бес% '::text))
        Rows Removed by Filter: 333333
Planning Time: 0.105 ms
Execution Time: 29.127 ms
(8 rows)
```

Таблицы и графики

Postgres					
Without index			With index		
concurrently requests	rate (throughput)	latency(avg), sec	concurrently requests	rate	latency(avg), sec
1	22	0,04	1	33	0,029
10	78	0,11	10	94	0,096
100	93	0,78	100	95	0,83
1000	93	7,2	1000	92	7,2



Выводы

1. Postgres не использует созданный индекс, хотя по идее должен.
2. До создания индекса запросы выполнялись на порядки быстрее чем с mysql(тоже без индекса)
3. Ввиду того, что postgres по какой-то причине не использует предложенный ему индекс, оценить время запроса с использованием индекса не представляется возможным.
4. Деградация, выражающаяся в увеличении latency наблюдается при throughput близким к 90 запросов в секунду.

Вопросы:

1. Как можно увеличить throughput сервера. Вероятно, какими-то настройками, но какими?
2. Как заставить postgres использовать индекс?
3. Почему mysql заметно проигрывает postgres без использования индекса?

Приложение 1 – лог запросов до добавления индекса

```
hey -n 10 -c 1 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'
```

Summary:
Total: 0.4513 secs
Slowest: 0.1454 secs
Fastest: 0.0315 secs
Average: 0.0451 secs
Requests/sec: 22.1604



Latency distribution:
10% in 0.0318 secs
25% in 0.0321 secs
50% in 0.0338 secs
75% in 0.0421 secs
90% in 0.1454 secs
0% in 0.0000 secs
0% in 0.0000 secs

Details (average, fastest, slowest):
DNS+ dialup: 0.0000 secs, 0.0315 secs, 0.1454 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0003 secs
req write: 0.0000 secs, 0.0000 secs, 0.0001 secs
resp wait: 0.0449 secs, 0.0313 secs, 0.1447 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0001 secs

Status code distribution:
[200] 10 responses

```
hey -n 100 -c 10 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'
```

Summary:
Total: 1.2690 secs
Slowest: 0.3524 secs
Fastest: 0.0298 secs
Average: 0.1126 secs
Requests/sec: 78.8030

Response time histogram:



Latency distribution:

10% in 0.0529 secs
25% in 0.0712 secs
50% in 0.0943 secs
75% in 0.1440 secs
90% in 0.2013 secs
95% in 0.2603 secs
99% in 0.3524 secs

Details (average, fastest, slowest):

DNS+ dialup: 0.0000 secs, 0.0298 secs, 0.3524 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0006 secs
req write: 0.0000 secs, 0.0000 secs, 0.0001 secs
resp wait: 0.1124 secs, 0.0297 secs, 0.3517 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0001 secs

Status code distribution:

[200] 100 responses

hey -n 1000 -c 100 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'

Summary:

Total: 10.7142 secs
Slowest: 9.0241 secs
Fastest: 0.0318 secs
Average: 0.7837 secs
Requests/sec: 93.3339

Response time histogram:



Latency distribution:

10% in 0.0814 secs
25% in 0.1269 secs
50% in 0.3089 secs
75% in 0.8841 secs
90% in 1.7725 secs
95% in 2.8852 secs
99% in 7.5123 secs

Details (average, fastest, slowest):

DNS+ dialup: 0.0003 secs, 0.0318 secs, 9.0241 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0026 secs
req write: 0.0001 secs, 0.0000 secs, 0.0058 secs
resp wait: 0.7831 secs, 0.0317 secs, 9.0182 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0002 secs

Status code distribution:

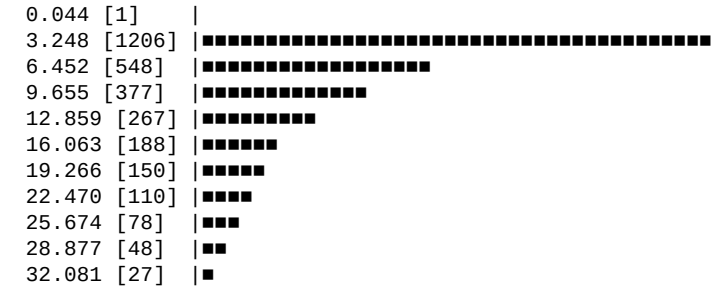
[200] 1000 responses

hey -n 3000 -c 1000 -t 0 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'

Summary:

Total: 32.2619 secs
Slowest: 32.0809 secs
Fastest: 0.0443 secs
Average: 7.2458 secs
Requests/sec: 92.9890

Response time histogram:



Latency distribution:

10% in 0.4512 secs
25% in 1.5501 secs
50% in 4.6728 secs
75% in 11.0035 secs
90% in 18.3749 secs
95% in 22.5326 secs
99% in 28.8461 secs

Details (average, fastest, slowest):

```
DNS+dialup: 0.0314 secs, 0.0443 secs, 32.0809 secs
DNS-lookup: 0.0217 secs, 0.0000 secs, 0.1522 secs
req write: 0.0030 secs, 0.0000 secs, 0.1451 secs
resp wait: 7.2083 secs, 0.0437 secs, 31.9330 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0083 secs
```

Status code distribution:

[200] 3000 responses

Приложение 2 – лог запросов после добавления индекса

```
hey -n 10 -c 1 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бес'
```

Summary:

```
Total:      0.2968 secs
Slowest:    0.0308 secs
Fastest:    0.0291 secs
Average:    0.0297 secs
Requests/sec: 33.6969
```

Response time histogram:

[illegible]

Latency distribution:

```
10% in 0.0293 secs
25% in 0.0293 secs
50% in 0.0297 secs
75% in 0.0301 secs
90% in 0.0308 secs
0% in 0.0000 secs
0% in 0.0000 secs
```

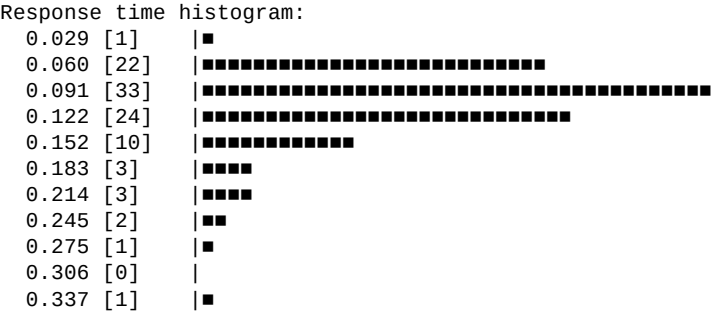
Details (average, fastest, slowest):

```
DNS+dialup: 0.0000 secs, 0.0291 secs, 0.0308 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0002 secs
req write: 0.0000 secs, 0.0000 secs, 0.0000 secs
resp wait: 0.0295 secs, 0.0290 secs, 0.0301 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0001 secs
```

Status code distribution:
[200] 10 responses

```
hey -n 100 -c 10 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'
```

Summary:
Total: 1.0527 secs
Slowest: 0.3370 secs
Fastest: 0.0294 secs
Average: 0.0964 secs
Requests/sec: 94.9944



Latency distribution:
10% in 0.0484 secs
25% in 0.0625 secs
50% in 0.0874 secs
75% in 0.1180 secs
90% in 0.1595 secs
95% in 0.2089 secs
99% in 0.3370 secs

Details (average, fastest, slowest):
DNS+ dialup: 0.0001 secs, 0.0294 secs, 0.3370 secs
DNS-lookup: 0.0000 secs, 0.0000 secs, 0.0004 secs
req write: 0.0000 secs, 0.0000 secs, 0.0003 secs
resp wait: 0.0962 secs, 0.0293 secs, 0.3369 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0001 secs

Status code distribution:
[200] 100 responses

```
hey -n 1000 -c 100 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бек'
```

Summary:
Total: 10.5230 secs
Slowest: 8.5217 secs

```
Fastest:      0.0294 secs
Average:     0.8302 secs
Requests/sec: 95.0295
```

Response time histogram:

[illegible]

Latency distribution:

```
10% in 0.0856 secs
25% in 0.1748 secs
50% in 0.4464 secs
75% in 1.0328 secs
90% in 2.0266 secs
95% in 3.0012 secs
99% in 5.4025 secs
```

Details (average, fastest, slowest):

```
DNS+ dialup: 0.0006 secs, 0.0294 secs, 8.5217 secs
DNS- lookup: 0.0000 secs, 0.0000 secs, 0.0058 secs
req write: 0.0004 secs, 0.0000 secs, 0.0162 secs
resp wait: 0.8290 secs, 0.0293 secs, 8.5095 secs
resp read: 0.0001 secs, 0.0000 secs, 0.0007 secs
```

Status code distribution:

[200] 1000 responses

```
hey -n 3000 -c 1000 -t 0 -m GET 'http://localhost:8800/user/search?first_name=Иван&second_name=Бес'
```

Summary:

```
Total:      32.5069  secs
Slowest:    31.8757  secs
Fastest:    0.0455  secs
Average:    7.2273  secs
Requests/sec: 92.2882
```

Response time histogram:

[illegible]

12.778	[249]		■■■■■■■■
15.961	[175]		■■■■■■■
19.144	[156]		■■■■■
22.327	[105]		■■■
25.510	[84]		■■■
28.693	[61]		■■
31.876	[29]		■

Latency distribution:

10%	in	0.4366	secs
25%	in	1.4828	secs
50%	in	4.5984	secs
75%	in	10.8024	secs
90%	in	18.6579	secs
95%	in	23.2156	secs
99%	in	28.6775	secs

Details (average, fastest, slowest):

DNS+ dialup:	0.0189	secs,	0.0455	secs,	31.8757	secs
DNS-lookup:	0.0056	secs,	0.0000	secs,	0.1119	secs
req write:	0.0020	secs,	0.0000	secs,	0.0695	secs
resp wait:	7.2017	secs,	0.0453	secs,	31.8354	secs
resp read:	0.0001	secs,	0.0000	secs,	0.0002	secs

Status code distribution:

[200]	3000	responses
-------	------	-----------